

Data Load Checklist

What the data owner does after the tables are created, in order, and what happens next. Reference data first, cases second, everything else stays empty until the models have run.

case_reference is not the first load. It has foreign keys to product and, if you added it, case_group, so those rows must exist first. Loading cases into an empty database will fail on the first row.

Decisions to make before loading AGREE THESE FIRST

DECISION	WHY IT MATTERS NOW
Product codes	Short, stable, 16 characters or fewer. They end up inside batch IDs such as B-1-STG9000, and the column holds 20.
Dataset version label	Stamped on every case row, e.g. DS-2026-10. If the sample is ever rebuilt, it gets a new label rather than an edit.
Case groups	At minimum a pilot group and a main group. The pilot sets the confidence threshold before full volume.
Pilot size	50 to 100 cases per product is enough to calibrate. These are SME-labelled blind, so keep it small.
Taxonomy version label	The first release, e.g. TAX-1.0. Every model run records which version it used.
Sampling strategy	How the cases are chosen: spread across skills, sub-products and case age, not just the most recent 1,000.

Load order FOREIGN KEYS DECIDE THIS

#	TABLE	ROWS	SOURCE
1	product	6	Typed by hand
2	taxonomy_version	1	Typed by hand
3	skill	~20	The agreed skill list
4	skill_definition	~20	One per skill for TAX-1.0, with inclusions and exclusions
5	case_group	2+	Pilot and main, if you added this table
6	case_reference	~6,000	Extract from the case system
7	sme	~8	The smaller reviewer group plus supervisors
8	sme_entitlement	~12	One row per SME per product they cover, dated

Everything else — model_run, model_output, auto_assessment, review_batch, review_queue, review_verdict, definition_change — stays empty. Those are filled by the model runs and the triage job. Do not pre-create batches or queue rows.

Steps IN ORDER

1 Create the schema

Data owner · half a day

Run the DDL against a test database first, then the real one. Confirm every table, constraint and index is created with no errors.

2 Load the reference data

Data owner · half a day

Items 1 to 5 above. Small enough to check by eye. The skill definitions are the important part: this is the wording that goes to both models, so it should be the approved text, not a placeholder.

3 Build and review the case sample

Data owner with product leads · 2 to 3 days

Extract candidate cases, apply the sampling strategy, and have a product lead glance at a sample per product before loading. Cases that are empty, duplicated or misfiled cost more to discover later.

4 Load case_reference

Data owner · half a day

Case number, product, title, description, URL, created date, sampling stratum, group, calibration flag, dataset version. Run the checks below before telling anyone it's done.

5 Load SMEs and entitlements

Data owner · 1 hour

Needs the names confirmed by the product Directors first. Set active_from to the date the review window opens, not today.

6 Freeze and record the sample

Gate

Data owner · 1 hour

Record the row counts per product and per group against the dataset version. From here the sample is read-only. If cases have to change, that is a new dataset version, not an update.

7 Hand over to data science

Data owner · 1 hour

Two things go out: the case numbers for the pilot group, and the full definition text for TAX-1.0. Confirm the data transfer has privacy sign-off before it leaves.

8 Pilot: run both models

First real output

Data science · days, not weeks

The primary model, and the second opinion run blind on the same cases. Two model_run rows, both loaded into model_output.

9 Pilot: SMEs label blind

SMEs · 2 to 3 days

SMEs review the pilot cases without seeing either model's answer. This is the only way to measure how accurate the second model is before you rely on it.

10 Set the confidence threshold

Gate

Project lead with data science · 1 day

Compare both models against the SME labels. Choose the threshold that gives an acceptable error rate in the auto-accepted pile. This decides what share of the main run is automated — the 80% is the result of this step, not an input to it.

11 Main run and triage

Data science, then the triage job

Both models over the main group, then the triage job writes `auto_assessment` and creates batches and queue rows for the cases that need a person. Batches start as Pending.

12 Open the first product

Project lead

Set one batch to Open. SMEs see work only from open batches, so products can be started one at a time rather than all six at once.

Checks after loading the cases BEFORE DECLARING IT DONE

```
-- 1. Count per product and group: expect the planned split
SELECT product_code, group_code, COUNT(*) AS cases,
       SUM(CASE WHEN is_golden_set = 1 THEN 1 ELSE 0 END) AS calibration
FROM validation.case_reference
GROUP BY product_code, group_code ORDER BY product_code, group_code;

-- 2. Duplicates: must return nothing
SELECT case_number, COUNT(*) FROM validation.case_reference
GROUP BY case_number HAVING COUNT(*) > 1;

-- 3. Missing content: cases the model cannot classify
SELECT COUNT(*) AS no_description FROM validation.case_reference
WHERE case_description IS NULL OR LEN(LTRIM(case_description)) < 50;

-- 4. Spread of case age: not all from the last month
SELECT MIN(case_created_date), MAX(case_created_date),
       COUNT(DISTINCT sampling_stratum) AS strata
FROM validation.case_reference;

-- 5. One dataset version only
SELECT dataset_version, COUNT(*) FROM validation.case_reference
GROUP BY dataset_version;
```

Empty or near-empty descriptions are worth removing now. They produce "can't tell" verdicts that look like model errors in the report and waste SME time.

Check the skill spread once the first model run lands. If rank 1 answers cluster on two or three skills, either the sample is skewed or something is wrong in the run.

What the owner should not do COMMON MISTAKES

Don't edit case_reference after freezing

Cycle comparison assumes the sample is fixed. Changes go in a new dataset version.

Don't pre-create batches or queue rows

The triage job creates them from the model output. Pre-created rows will collide with it.

Don't load the full sample before the pilot

Loading is fine; *running* the full sample before the threshold is set wastes a model run.

Don't reuse calibration cases to report final accuracy

Cases used to set the threshold will flatter the result.

Don't let placeholder definitions reach data science

The wording is the thing being tested as much as the model is.