

Here's a naming standard and a map of where each dimension lives.

Naming standard

Format: PREFIX-[wave]-[cycle]-[scope]-[sequence], uppercase, hyphen separated, segments always in that order, sequences zero-padded to a fixed width.

Entity	Pattern	Example	Reads as
Wave (group)	W{n}-YYYY-MM	W1-2026-10	Wave 1, extracted Oct 2026
Product	up to 8 chars, no hyphens	STG9000	Storage array 9000
Model run	R-{wave}-C{cycle}-{P S}	R-W1-C2-P	Wave 1, cycle 2, primary
Batch	B-{wave}-C{cycle}-{product}	B-W1-C2-STG9000	That product's share of that run
Queue item	Q-{wave}-C{cycle}-{nnnnnn}	Q-W1-C2-000148	
Verdict	V-{nnnnnnnn}	V-0000912	
Definition change	DC-{nnnn}	DC-0014	
Definition package	PKG-{nnnn}	PKG-0007	
Case	source case number, unchanged	CS-0400001	

Five rules behind it:

- 1 Encode only what someone needs to read at a glance** in an exported spreadsheet. Wave, cycle, product and role qualify. Slot, status and priority don't — they're columns.
- 2 Never parse an ID in code.** They're human aids. Every join goes through the foreign key. If anything in the application splits a batch ID to get a product, that's a defect waiting for the first product rename.
- 3 Nothing mutable in an identifier.** Status and priority change; the ID must not.
- 4 Fixed-width sequences**, so sorting alphabetically sorts correctly.
- 5 Product codes capped at 8 characters.** B-W1-C2- is 8, and the column holds 20. A longer product code silently breaks batch IDs.

Verdicts, changes and packages get bare sequences because they're always reached by join, never read in isolation.

Where each dimension lives

Where at you want	Table.column	Reached from the queue by
----------------------------	--------------	---------------------------

Wa	case_group.group_code	review_queue.case_num
ve		ber →
		case_reference.group_code
	also model_run.group_code	review_queue.run_id →
		model_run.group_code
Pro	product.product_code	review_queue.batch_id →
duct		review_batch.product_code
Cyc	model_run.cycle_number	review_queue.run_id →
le		model_run
Run	model_run.run_role,	same
role	paired_run_id	
Bat	review_batch.batch_id, status,	direct FK
ch	target_size	
Que	review_queue.queue_id	—
ue		
item		
Slot	review_queue.review_slot	1 = normal, 2 = second reviewer
Wh		
y	review_queue.routing_reason	
que		
ued		
Blin	review_queue.is_blind	
d		
Cali		
brat	case_reference.is_golden_set	
ion		
case		
Rev	review_queue.assigned_sme	
iewe	(claimed),	
r	review_verdict.viewer(actual)	
Ver	review_verdict.verdict_id,	queue_id
dict	sided_with, error_cat	
Tria	auto_assessment.outcome,	(case_number,
ge	skill_agree, object_agree,	primary_run_id)
outc	defect_agree, rule_version	
ome		

Tax
ono
my model_run.taxonomy_version a label only
labe
l
Act
ual v_run_definitions over
wor definition_package run_id
ding
Cha definition_change.change_id →
nge skill_definition_history.change
line _id
age _id

Two deliberate absences: the queue has no product column, and `case_reference` has no cycle column. Both are derived, and duplicating either is how they drift.

One query that returns every dimension

```

SELECT g.group_code           AS wave,
       b.product_code         AS product,
       r.cycle_number         AS cycle,
       r.run_role,
       b.batch_id,
       q.queue_id, q.review_slot, q.routing_reason,
       q.is_blind,
       c.is_golden_set         AS is_calibration,
       a.outcome               AS triage_outcome,
       v.verdict_id, vReviewer, v.sided_with,
       v.error_cat
FROM validation.review_queue q
JOIN validation.review_batch  b ON b.batch_id =
q.batch_id AND b.run_id = q.run_id
JOIN validation.model_run     r ON r.run_id = q.run_id
JOIN validation.case_reference c ON c.case_number =
q.case_number
JOIN validation.case_group    g ON g.group_code =
c.group_code
LEFT JOIN validation.auto_assessment a
      ON a.case_number = q.case_number AND
a.primary_run_id = q.run_id
LEFT JOIN validation.review_verdict v
      ON v.queue_id = q.queue_id AND
v.supersedes_verdict_id IS NULL;

```

Worth making this a view — `v_review_fact` or similar — so reporting and the

supervisor screen aren't each rebuilding the join chain.

One inconsistency to fix now

The controlled vocabularies don't currently agree on case: statuses are Title Case (`Queued`, `In Review`), error categories are lower case with spaces (`model error`), routing reasons are lower case (`llm disagreement`). It works, but every query that filters on one has to remember which convention applies, and someone will eventually write `WHERE status = 'queued'` and get nothing back.

Standardise on lower snake case for all of them — `queued`, `in_review`, `model_error`, `llm_disagreement` — while the tables are still empty. Changing check constraints now costs minutes; after the first cycle it means a data migration and rewriting every report.